

Duncan530Week10FinalProject

August 8, 2024

```
[40]: import pandas as pd
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from scipy import stats
from scipy.stats import skew, kurtosis
# Load the dataset
url = "https://github.com/taylorduncan/DSC530/raw/main/
↳Most%20Streamed%20Spotify%20Songs%202024.csv"

try:
    df = pd.read_csv(url, encoding='ISO-8859-1')
except UnicodeDecodeError as e:
    print(f"UnicodeDecodeError: {e}")

# Display the first few rows of the dataset
df.head()

# Display the dataset columns and basic statistics
df.info()
df.describe(include='all')
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4600 entries, 0 to 4599
```

```
Data columns (total 29 columns):
```

#	Column	Non-Null Count	Dtype
0	Track	4600 non-null	object
1	Album Name	4600 non-null	object
2	Artist	4595 non-null	object
3	Release Date	4600 non-null	object
4	ISRC	4600 non-null	object
5	All Time Rank	4600 non-null	object
6	Track Score	4600 non-null	float64
7	Spotify Streams	4487 non-null	object
8	Spotify Playlist Count	4530 non-null	object
9	Spotify Playlist Reach	4528 non-null	object

10	Spotify Popularity	3796	non-null	float64
11	YouTube Views	4292	non-null	object
12	YouTube Likes	4285	non-null	object
13	TikTok Posts	3427	non-null	object
14	TikTok Likes	3620	non-null	object
15	TikTok Views	3619	non-null	object
16	YouTube Playlist Reach	3591	non-null	object
17	Apple Music Playlist Count	4039	non-null	float64
18	AirPlay Spins	4102	non-null	object
19	SiriusXM Spins	2477	non-null	object
20	Deezer Playlist Count	3679	non-null	float64
21	Deezer Playlist Reach	3672	non-null	object
22	Amazon Playlist Count	3545	non-null	float64
23	Pandora Streams	3494	non-null	object
24	Pandora Track Stations	3332	non-null	object
25	Soundcloud Streams	1267	non-null	object
26	Shazam Counts	4023	non-null	object
27	TIDAL Popularity	0	non-null	float64
28	Explicit Track	4600	non-null	int64

dtypes: float64(6), int64(1), object(22)

memory usage: 1.0+ MB

[40]:

	Track	Album Name	Artist	Release Date	\
count	4600	4600	4595	4600	
unique	4370	4005	1999	1562	
top	Danza Kuduro - Cover	Un Verano Sin Ti	Drake	1/1/2012	
freq	13	20	63	38	
mean	NaN	NaN	NaN	NaN	
std	NaN	NaN	NaN	NaN	
min	NaN	NaN	NaN	NaN	
25%	NaN	NaN	NaN	NaN	
50%	NaN	NaN	NaN	NaN	
75%	NaN	NaN	NaN	NaN	
max	NaN	NaN	NaN	NaN	

	ISRC	All Time Rank	Track Score	Spotify Streams	\
count	4600	4600	4600.000000	4487	
unique	4598	4577	NaN	4425	
top	USWL11700269	3,441	NaN	1,655,575,417	
freq	2	2	NaN	4	
mean	NaN	NaN	41.844043	NaN	
std	NaN	NaN	38.543766	NaN	
min	NaN	NaN	19.400000	NaN	
25%	NaN	NaN	23.300000	NaN	
50%	NaN	NaN	29.900000	NaN	
75%	NaN	NaN	44.425000	NaN	
max	NaN	NaN	725.400000	NaN	

	Spotify Playlist Count	Spotify Playlist Reach	...	SiriusXM Spins	\
count	4530	4528	...	2477	
unique	4207	4478	...	689	
top	1	3	...	1	
freq	46	8	...	54	
mean	NaN	NaN	...	NaN	
std	NaN	NaN	...	NaN	
min	NaN	NaN	...	NaN	
25%	NaN	NaN	...	NaN	
50%	NaN	NaN	...	NaN	
75%	NaN	NaN	...	NaN	
max	NaN	NaN	...	NaN	

	Deezer Playlist Count	Deezer Playlist Reach	Amazon Playlist Count	\
count	3679.000000	3672	3545.000000	
unique	NaN	3558	NaN	
top	NaN	1,097	NaN	
freq	NaN	17	NaN	
mean	32.310954	NaN	25.348942	
std	54.274538	NaN	25.989826	
min	1.000000	NaN	1.000000	
25%	5.000000	NaN	8.000000	
50%	15.000000	NaN	17.000000	
75%	37.000000	NaN	34.000000	
max	632.000000	NaN	210.000000	

	Pandora Streams	Pandora Track Stations	Soundcloud Streams	\
count	3494	3332	1267	
unique	3491	2975	1265	
top	56,972,562	9	27	
freq	2	6	2	
mean	NaN	NaN	NaN	
std	NaN	NaN	NaN	
min	NaN	NaN	NaN	
25%	NaN	NaN	NaN	
50%	NaN	NaN	NaN	
75%	NaN	NaN	NaN	
max	NaN	NaN	NaN	

	Shazam Counts	TIDAL Popularity	Explicit Track
count	4023	0.0	4600.000000
unique	4002	NaN	NaN
top	1	NaN	NaN
freq	5	NaN	NaN
mean	NaN	NaN	0.358913
std	NaN	NaN	0.479734

min	NaN	NaN	0.000000
25%	NaN	NaN	0.000000
50%	NaN	NaN	0.000000
75%	NaN	NaN	1.000000
max	NaN	NaN	1.000000

[11 rows x 29 columns]

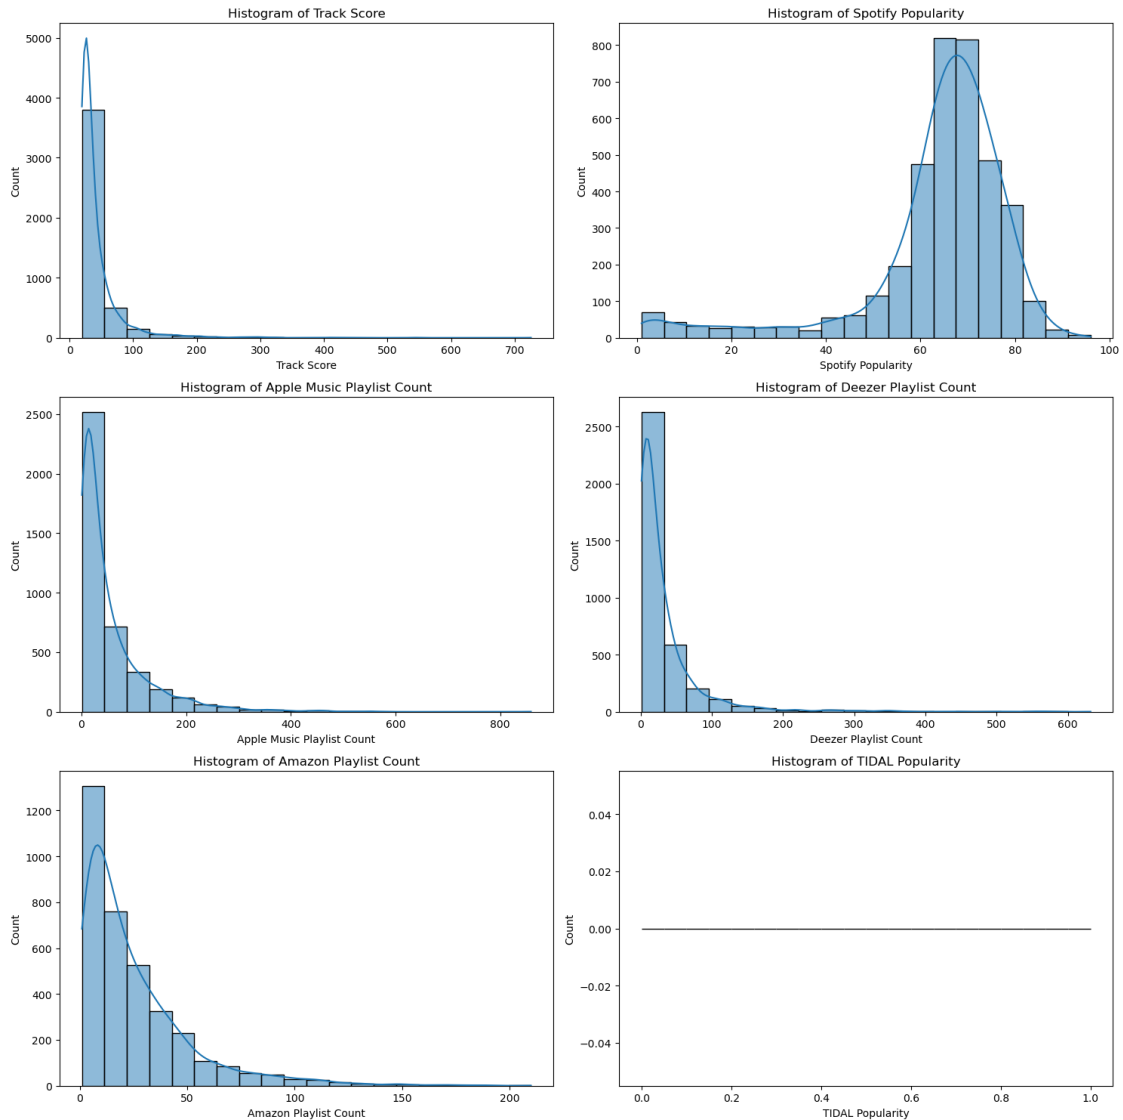
```
[ ]: '''
Describe these variables:
The dataset contains the following variables:
1.      Track Score: Likely a numerical value indicating the overall score or
      ↳ ranking of the track based on certain criteria (e.g., popularity, user
      ↳ ratings).
2.      Spotify Popularity: A numerical score representing how popular the
      ↳ track is on Spotify, possibly on a scale from 0 to 100.
3.      Apple Music Playlist Count: The number of playlists that feature the
      ↳ track on Apple Music.
4.      Deezer Playlist Count: The number of playlists that feature the track
      ↳ on Deezer.
5.      Amazon Playlist Count: The number of playlists that feature the track
      ↳ on Amazon Music.
6.      TIDAL Popularity: A numerical score representing the track's
      ↳ popularity on TIDAL, similar to Spotify Popularity.
'''
```

```
[34]: #Plot histograms for each of the numerical variables and calculate descriptive
      ↳ statistics.

fig, axes = plt.subplots(3, 2, figsize=(15, 15))
axes = axes.flatten()

for idx, column in enumerate(['Track Score', 'Spotify Popularity', 'Apple Music
      ↳ Playlist Count', 'Deezer Playlist Count', 'Amazon Playlist Count', 'TIDAL
      ↳ Popularity']):
    sns.histplot(df[column], bins=20, kde=True, ax=axes[idx])
    axes[idx].set_title(f'Histogram of {column}')

plt.tight_layout()
plt.show()
```



[77]: *#Descriptive Statistics: Calculate the mean, mode, spread, and tails (skewness and kurtosis) for each variable.*

```
# Identify columns with all NaN values
all_nan_columns = df.columns[df.isna().all()].tolist()

# Drop columns with all NaN values if needed
df = df.drop(columns=all_nan_columns)

# Select only numeric columns
numeric_df = df.select_dtypes(include=[float, int])
# Fill missing values for numeric columns
df_filled = numeric_df.fillna(numeric_df.mean())
```

```

from scipy.stats import skew, kurtosis
# Initialize dictionary to store descriptive statistics
descriptive_stats = {}
# Iterate over each column to calculate descriptive statistics
for column in numeric_df.columns:
    try:
        # Calculate statistics
        mean = df_filled[column].mean()
        mode = df_filled[column].mode()
        mode_value = mode.iloc[0] if not mode.empty else None
        spread = df_filled[column].std()
        skewness = skew(df_filled[column].dropna())
        kurt = kurtosis(df_filled[column].dropna())

        # Store statistics in dictionary
        descriptive_stats[column] = {
            'Mean': mean,
            'Mode': mode_value,
            'Spread (Std Dev)': spread,
            'Skewness': skewness,
            'Kurtosis': kurt
        }
    except KeyError:
        descriptive_stats[column] = 'Column not found'
    except Exception as e:
        descriptive_stats[column] = str(e)

# Display descriptive statistics
descriptive_stats

```

```

[77]: {'Track Score': {'Mean': 41.84404347826101,
    'Mode': 21.7,
    'Spread (Std Dev)': 38.54376574715532,
    'Skewness': 5.686368434584364,
    'Kurtosis': 53.87462816507168},
    'Spotify Popularity': {'Mean': 63.501580611169146,
    'Mode': 63.50158061116965,
    'Spread (Std Dev)': 14.7036610951589,
    'Skewness': -2.257955847370638,
    'Kurtosis': 6.286506435872026},
    'Apple Music Playlist Count': {'Mean': 54.60311958405621,
    'Mode': 54.603119584055456,
    'Spread (Std Dev)': 67.10252951585504,
    'Skewness': 3.0810885211454995,
    'Kurtosis': 14.830274224929838},
    'Deezer Playlist Count': {'Mean': 32.31095406360473,
    'Mode': 32.31095406360424,

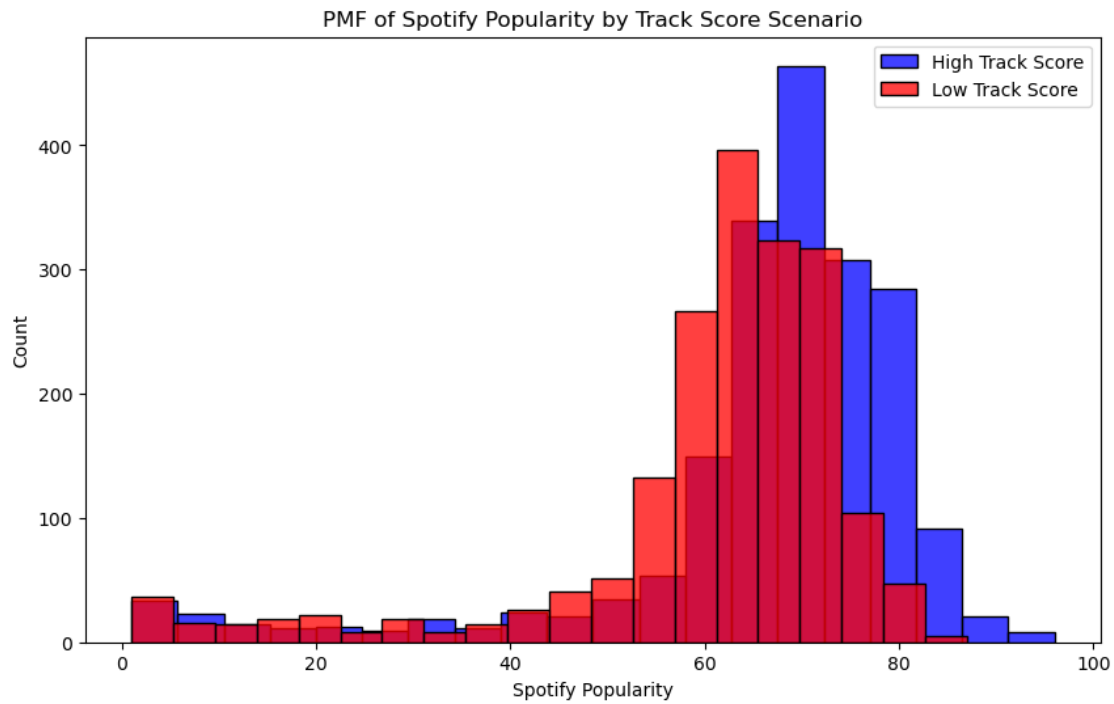
```

```
'Spread (Std Dev)': 48.53670525065182,
'Skewness': 5.358580903488518,
'Kurtosis': 41.3533812484804},
'Amazon Playlist Count': {'Mean': 25.348942172073333,
'Mode': 25.348942172073343,
'Spread (Std Dev)': 22.814897297845793,
'Skewness': 2.4890521710106395,
'Kurtosis': 9.32674440860418},
'Explicit Track': {'Mean': 0.35891304347826086,
'Mode': 0,
'Spread (Std Dev)': 0.47973378254293886,
'Skewness': 0.5882524837561582,
'Kurtosis': -1.653959015354711}}
```

[51]: *#Probability Mass Function (PMF): Compare two scenarios using the PMF. For*
↪instance, compare the distribution of Spotify Popularity when the Track
↪Score is above and below its median.

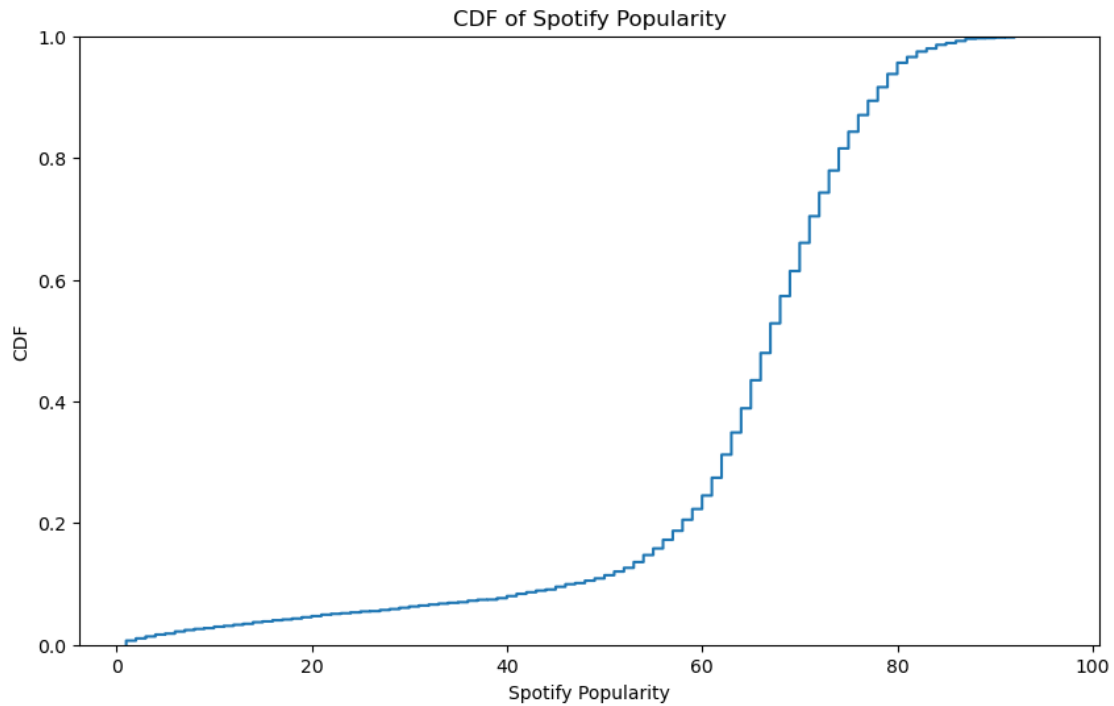
```
# Create two scenarios based on Track Score median
median_score = df['Track Score'].median()
high_score = df[df['Track Score'] > median_score]['Spotify Popularity']
low_score = df[df['Track Score'] <= median_score]['Spotify Popularity']

# Plot PMF
fig, ax = plt.subplots(figsize=(10, 6))
sns.histplot(high_score, bins=20, kde=False, label='High Track Score',
↪color='blue', ax=ax)
sns.histplot(low_score, bins=20, kde=False, label='Low Track Score',
↪color='red', ax=ax)
ax.set_title('PMF of Spotify Popularity by Track Score Scenario')
ax.legend()
plt.show()
```



[52]: *#Cumulative Distribution Function (CDF): Create a CDF for Spotify Popularity.*

```
# Plot CDF
plt.figure(figsize=(10, 6))
sns.ecdfplot(df['Spotify Popularity'])
plt.title('CDF of Spotify Popularity')
plt.xlabel('Spotify Popularity')
plt.ylabel('CDF')
plt.show()
```

[55]: *#Analytical Distribution: Fit a distribution (e.g., Normal) and compare it to*
↳ the actual data.

```
# Check for non-finite values in 'Spotify Popularity'
non_finite_values = df['Spotify Popularity'].isna().sum() + np.
↳ isinf(df['Spotify Popularity']).sum()
print(f'Number of non-finite values: {non_finite_values}')

# Drop non-finite values
spotify_popularity_clean = df['Spotify Popularity'].dropna()
spotify_popularity_clean = spotify_popularity_clean[np.
↳ isfinite(spotify_popularity_clean)]

from scipy.stats import norm
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Fit a normal distribution to the cleaned data
mu, std = norm.fit(spotify_popularity_clean)

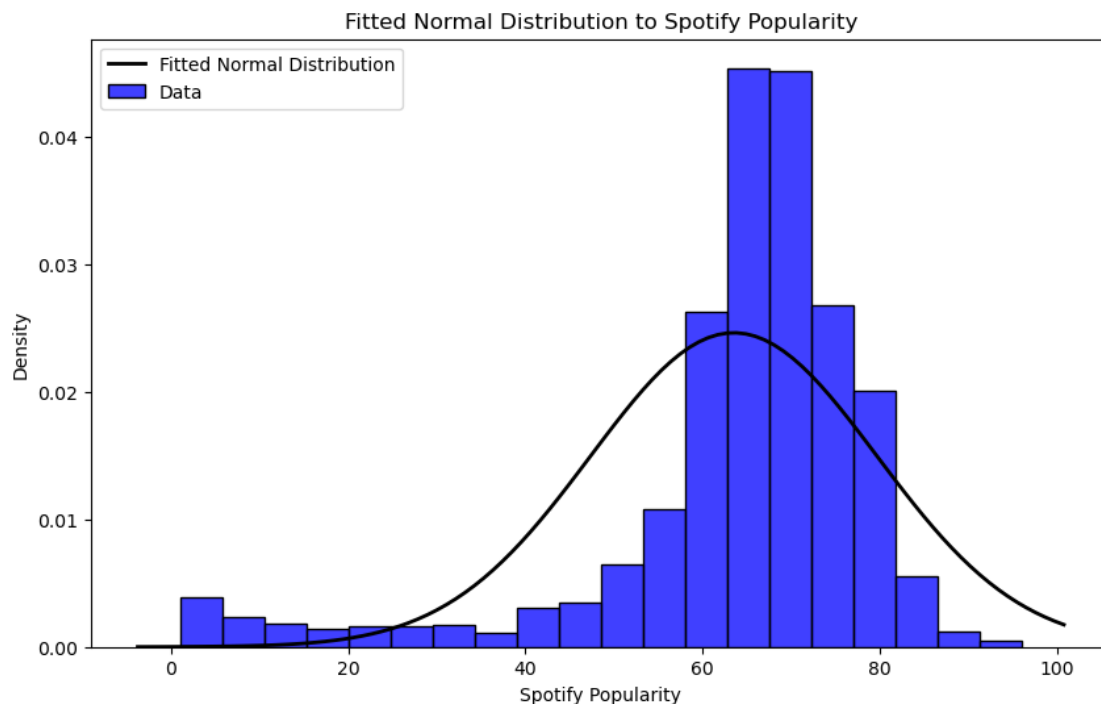
# Plot the distribution
plt.figure(figsize=(10, 6))
```

```

sns.histplot(spotify_popularity_clean, bins=20, kde=False, label='Data',
             color='blue', stat='density')
xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
p = norm.pdf(x, mu, std)
plt.plot(x, p, 'k', linewidth=2, label='Fitted Normal Distribution')
plt.title('Fitted Normal Distribution to Spotify Popularity')
plt.xlabel('Spotify Popularity')
plt.ylabel('Density')
plt.legend()
plt.show()

```

Number of non-finite values: 804



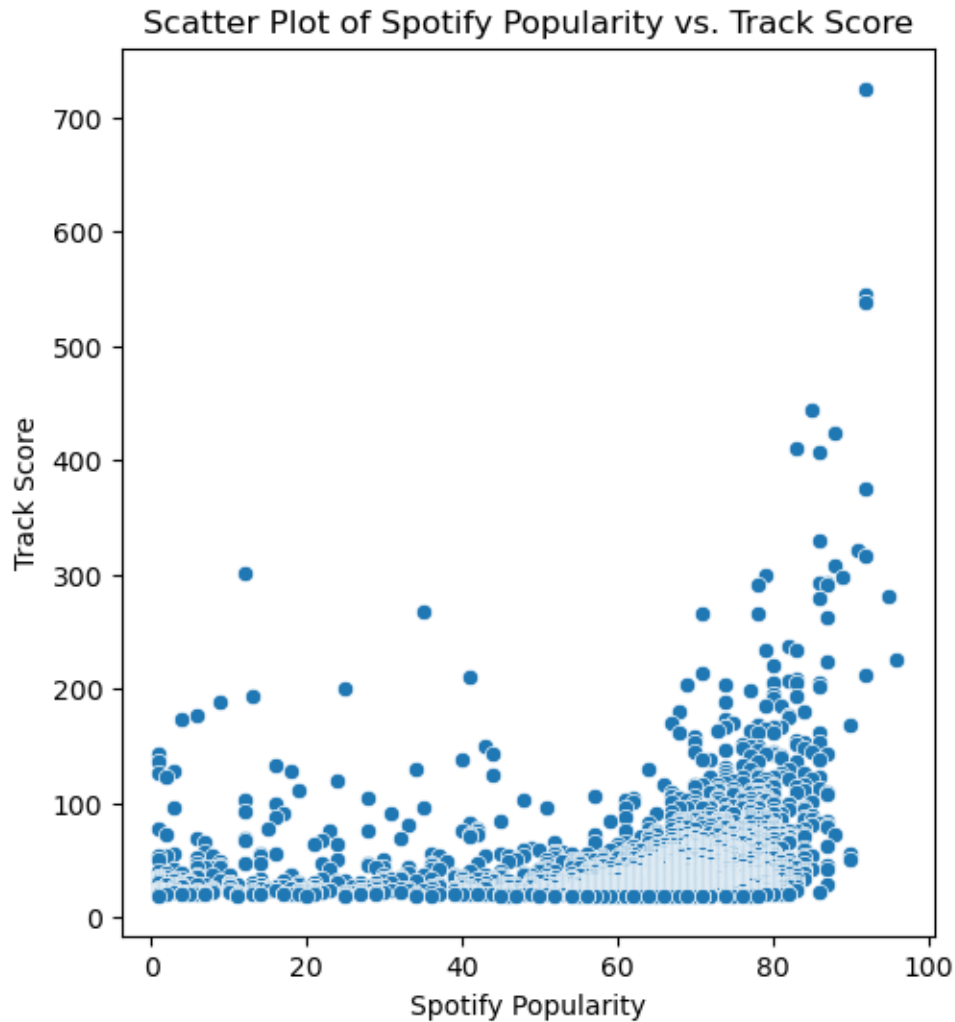
[74]: *#Scatter Plots and Correlation Analysis: Compare two variables with scatter plots and compute correlations.*

```

# Scatter plot between Spotify Popularity and Track Score
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
sns.scatterplot(data=df, x='Spotify Popularity', y='Track Score')
plt.title('Scatter Plot of Spotify Popularity vs. Track Score')
plt.xlabel('Spotify Popularity')
plt.ylabel('Track Score')

```

```
plt.show()
# Compute correlation
correlation = df[['Spotify Popularity', 'Track Score']].corr().iloc[0, 1]
correlation
```

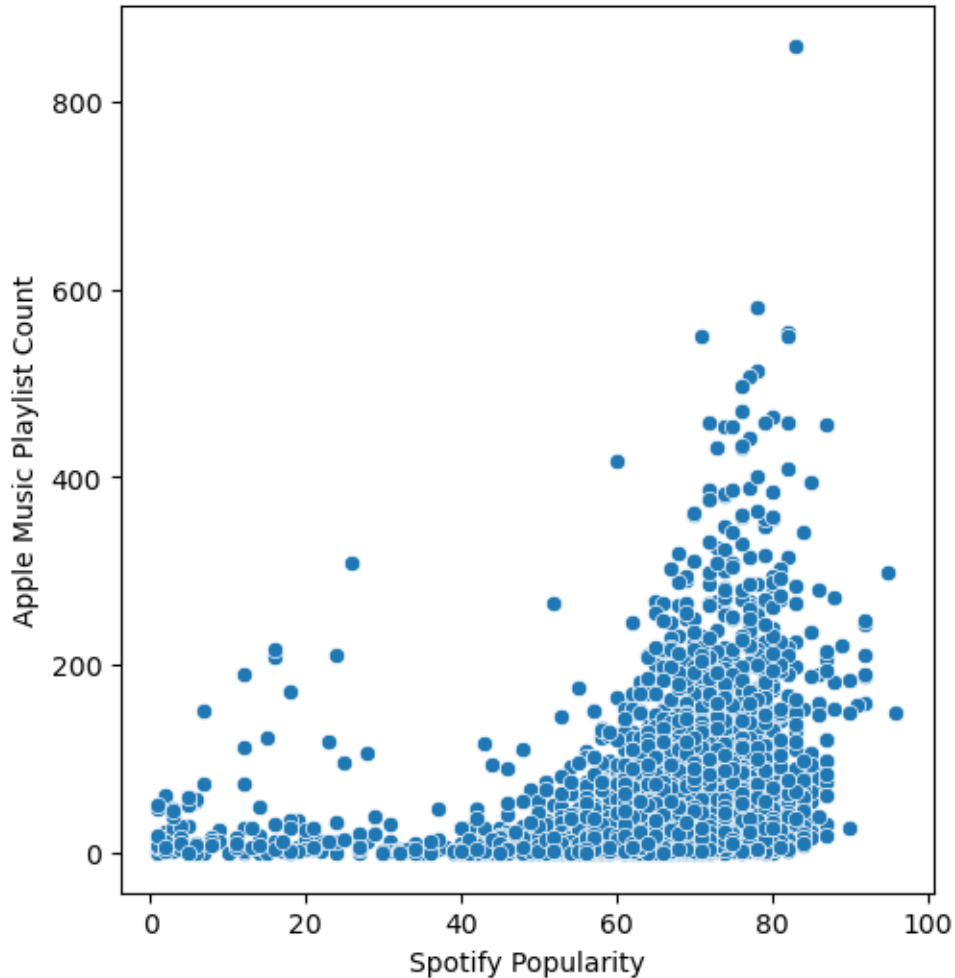


[74]: 0.19280775654813445

```
[75]: # Scatter plot between Spotify Popularity and Track Score
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 2)
sns.scatterplot(data=df, x='Spotify Popularity', y='Apple Music Playlist Count')
plt.title('Scatter Plot of Spotify Popularity vs. Apple Music Playlist Count')
plt.xlabel('Spotify Popularity')
plt.ylabel('Apple Music Playlist Count')
plt.show()
```

```
# Compute correlation
correlation = df[['Spotify Popularity', 'Apple Music Playlist Count']].corr().
    ↳iloc[0, 1]
correlation
```

Scatter Plot of Spotify Popularity vs. Apple Music Playlist Count



[75]: 0.2884300282744505

[57]: *#Hypothesis Testing: Perform a hypothesis test, for example, testing if the*
↳mean Spotify Popularity differs from a known value.

```
from scipy.stats import ttest_1samp

# Hypothesis test: mean Spotify Popularity differs from 50
test_statistic, p_value = ttest_1samp(df['Spotify Popularity'].dropna(), 50)
test_statistic, p_value
```

[57]: (51.392106212822206, 0.0)

[69]: *#Regression Analysis: Perform regression analysis to predict one variable based on others. For instance, predicting Spotify Popularity from Track Score and playlist counts.*

```
import statsmodels.api as sm

# Option 1: Drop rows with any missing values
df_cleaned = df.dropna()

# Option 2: Fill missing values with the mean of each column
df_filled = df.fillna(df.mean())

df_cleaned.replace([float('inf'), -float('inf')], np.nan, inplace=True)
df_cleaned.dropna(inplace=True) # Remove rows with any NaN values

# Prepare data for regression with cleaned dataset
X = df_cleaned[['Track Score', 'Apple Music Playlist Count', 'Deezer Playlist_
↳Count', 'Amazon Playlist Count']]
X = sm.add_constant(X) # Adds a constant term to the predictors
y = df_cleaned['Spotify Popularity']

# Fit regression model
model = sm.OLS(y, X).fit()
print(model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          Spotify Popularity    R-squared:                0.172
Model:                  OLS                  Adj. R-squared:           0.166
Method:                 Least Squares        F-statistic:             29.13
Date:                  Thu, 08 Aug 2024      Prob (F-statistic):      5.05e-22
Time:                  22:49:30              Log-Likelihood:         -2111.5
No. Observations:      565                  AIC:                    4233.
Df Residuals:          560                  BIC:                    4255.
Df Model:               4
Covariance Type:       nonrobust
=====
=====
```

		coef	std err	t	P> t
[0.025	0.975]				

const		64.1015	0.698	91.889	0.000
62.731	65.472				
Track Score		0.0436	0.009	5.078	0.000

0.027	0.060				
Apple Music Playlist Count	0.0199	0.010	2.066	0.039	
0.001	0.039				
Deezer Playlist Count	-0.0072	0.013	-0.554	0.580	
-0.033	0.018				
Amazon Playlist Count	0.0411	0.021	1.954	0.051	
-0.000	0.082				
=====					
Omnibus:	394.388	Durbin-Watson:		1.962	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		5421.998	
Skew:	-2.953	Prob(JB):		0.00	
Kurtosis:	16.980	Cond. No.		263.	
=====					